

Information Retrieval and Data Mining (COMP0084)

Coursework 1

Task definition

An information retrieval model is an essential component for many applications such as search, question answering, and recommendation. In this coursework you are going to develop information retrieval models that solve the problem of passage retrieval, i.e. given a query, return a ranked list of short texts (passages). More specifically, you are going to build a passage re-ranking system: given a candidate list of passages for a query, you should re-rank these candidate passages based on an information retrieval model.

qid	pid	query	passage
523270	2818345	toyota of plano plano tx	DART's Red Line runs along North Central Expre...
527433	1537731	types of dysarthria from cerebral palsy	In some cases, further testing will also be ab...
1113437	5194230	what is physical description of spruce	Source: "U.S. Rehab Aide Job Description. Reha...
833860	5043973	what is the most popular food in switzerland	The national currency in Switzerland is the Sw...
1056204	2328990	who was the first steam boat operator	a relatively small usually open craft of a siz...

Figure 1: Sample rows from the `candidate-passages-top1000.tsv` file.

Data

Data can be downloaded from [this online repository](#).¹ The data set consists of 3 files:

- `test-queries.tsv` is a tab separated file, where each row contains a test query identifier (`qid`) and the actual query text.
- `passage-collection.txt` is a collection of passages, one per row.
- `candidate-passages-top1000.tsv` is a tab separated file with an initial selection of at most 1000 passages for each of the queries in `test-queries.tsv`. The format of this file is `<qid pid query passage>`, where `pid` is the identifier of the passage retrieved, `query` is the query text, and `passage` is the passage text (all tab separated). The passages contained in this file are the same as the ones in `passage-collection.txt`. However, there might be

¹Link to the data, <https://www.dropbox.com/s/z15xt8og1a7bt3q/coursework-1-data.zip?dl=0>

some repetitions, i.e. the same passage could have been returned for more than 1 queries. Figure 1 shows some sample rows from that file.

It is important that you do **not** change the above filenames, i.e. your source code must use the above filenames, otherwise we might not be able to assess your code automatically.

Coursework tasks and deliverables

Please read carefully as all deliverables (**D1** to **D12**) are described within the tasks. If a deliverable requests an answer in a specific format, follow these guidelines carefully as marking may be automated (incorrect format will be penalised and might result to 0 marks). Although we strongly suggest using Python, you could also use Java. Please make one consistent choice for this coursework. Do not submit your source code as a Jupyter Notebook or similar. Do not use external functions or libraries that can solve (end-to-end) the tasks of building an inverted index, TF-IDF, the vector space model, BM25, or any of the query likelihood language models. You are not allowed to reuse any code that is available from someone or somewhere else (i.e. write your own code). Use only unigram (1-gram) text representations to solve this coursework's tasks.² Write your report using the ACL L^AT_EX template,³ and submit 1 PDF file named **report.pdf**. All plots in your report must be in vector format (e.g. PDF). Your report should not exceed 6 pages. In total, your submission should be made of 4 source code files, 5 output CSV files, and 1 PDF file with the report. Do not deploy any internal directory structure in your submission, i.e. just submit the aforementioned 10 files.

Task 1 – Text statistics (30 marks). Use the data in `passage-collection.txt` for this task. Extract terms (1-grams) from the raw text. In doing so, you can also perform basic text preprocessing steps. However, you can also choose not to. In any case, you should not remove stop words for Task 1, but you can do so in later tasks of the coursework.

D1 Describe and justify your preprocessing choice(s), if any, and report the size of the identified index of terms (vocabulary). Then, implement a function that counts the number of occurrences of terms in the provided data set, plot their probability of occurrence (normalised frequency) against their frequency ranking, and qualitatively justify that these terms follow Zipf's law. As a reminder, Zipf's law for text sets $s = 1$ in the Zipfian distribution defined by

$$f(k; s, N) = \frac{k^{-s}}{\sum_{i=1}^N i^{-s}}, \quad (1)$$

where $f(\cdot)$ denotes the normalised frequency of a term, k denotes the term's frequency rank in our corpus (with $k = 1$ being the highest rank), N is the number of terms in our vocabulary, and s is a distribution parameter. How does your empirical distribution compare to the actual Zipf's law distribution? Use Eq. 1 to explain where their difference is coming from, and also compare the two in a log-log plot.

D2 Submit your source code for Task 1 as a single Python or Java file titled **task1.py** or **task1.java**, respectively.

² n -gram representations with $n > 1$ can improve performance, but are out-of-scope for this coursework.

³The ACL L^AT_EX template is available as an OverLeaf project, [overleaf.com/read/crtcgwxzjskr](https://www.overleaf.com/read/crtcgwxzjskr).

Task 2 – Inverted index (15 marks). Use `candidate-passages-top1000.tsv` for this task (unique instances of column pairs `pid` and `passage`). Using the vocabulary of terms identified in Task 1 (you might also consider removing stop words), build an inverted index for the collection so that you can retrieve passages in an efficient way.

D3 Provide a description of your approach to generating an inverted index, report the information that you decided to store in it, and justify this choice.

Hint: *Implementing Tasks 3 and 4 should inform this choice.*

D4 Submit the source code for Task 2 as a single Python or Java file titled `task2.py` or `task2.java`, respectively.

Task 3 – Retrieval models (35 marks). Use `test-queries.tsv` and `candidate-passages-top1000.tsv` for this task.

D5 Extract the TF-IDF vector representations of the passages using the inverted index you have constructed. Using the passages' IDF representation, extract the TF-IDF representation of the queries. Use a basic vector space model with TF-IDF and cosine similarity to retrieve at most 100 passages from within the 1000 candidate passages for each query (some queries have fewer candidate passages). Store the outcomes in a file named `tfidf.csv`. Each row of `tfidf.csv` must have the following format:

`qid,pid,score`

where `qid` denotes the query's identifier, `pid` denotes the passage identifier, and `score` is their cosine similarity score. Note that there should be no spaces between commas. Strictly report the identifiers from the raw data (do not replace them with your own). `qid,pid` pairs should be ranked by similarity score in descending order, i.e. starting with the highest similarity score. You should first report the top `qid,pid` pairs for one `qid`, then move on to the next, following the same query order as in the `test-queries.tsv` file. The same file format should be used for reporting the outputs of all the retrieval or language models that you will implement as part of this coursework. The last column should report the corresponding similarity/relevance score of each model.

D6 Use the inverted index to also implement BM25⁴ while setting $k_1 = 1.2$, $k_2 = 100$, and $b = 0.75$. Retrieve at most 100 passages from within the 1000 candidate passages for each test query. Store the outcomes in a file named `bm25.csv`.

D7 Submit the source code for Task 3 as a single Python or Java file titled `task3.py` or `task3.java`, respectively.

Task 4 – Query likelihood language models (20 marks). Use `test-queries.tsv` and `candidate-passages-top1000.tsv` for this task. Implement the query likelihood language model with (a) Laplace smoothing, (b) Lidstone correction with $\epsilon = 0.1$, and (c) Dirichlet smoothing with $\mu = 50$, and retrieve 100 passages from within the 1000 candidate passages for each test query.

D8-10 Store the respective outcomes in the files `laplace.csv`, `lidstone.csv`, and `dirichlet.csv`. In these files, the column `score` should report the natural logarithm of the probability scores.

⁴BM25 is explained in detail in [this paper](#).

- D11** Which language model do you expect to work better? Which ones are expected to be more similar and why? Comment on the value of $\epsilon = 0.1$ in the Lidstone correction – is this a good choice, and why? If we set $\mu = 5000$ in Dirichlet smoothing, would this be a more appropriate value and why?
- D12** Submit the source code for Task 4 as a single Python or Java file titled `task4.py` or `task4.java`, respectively.

Acknowledgements

I would like to thank Omer Kirnap and Sahan Bulathwela for developing the original version of this coursework.